

Workshop: Intro to Bayesian Learning

Lecture 5 - Introduction to Gibbs sampling, MCMC and HMC

Mattias Villani

**Department of Statistics
Stockholm University**



Overview

- Gibbs sampling
- The Metropolis-Hastings algorithm
- Hamiltonian Monte Carlo

Monte Carlo sampling

- If $\theta^{(1)}, \dots, \theta^{(m)}$ is an **iid sequence** from $p(\theta|\mathbf{y})$, then

$$\bar{\theta} = \frac{1}{m} \sum_{i=1}^m \theta^{(i)} \rightarrow \mathbb{E}(\theta|\mathbf{y})$$

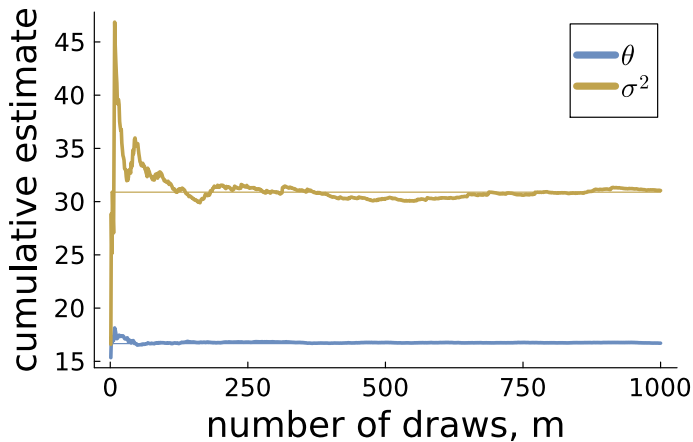
$$\bar{g}(\theta) = \frac{1}{m} \sum_{i=1}^m g(\theta^{(i)}) \rightarrow \mathbb{E}[g(\theta)|\mathbf{y}]$$

for some function $g(\theta)$ of interest.

- **Central limit theorem**

$$\bar{\theta}_{1:m} \stackrel{\text{appr}}{\sim} N\left(\mathbb{E}(\theta|\mathbf{y}), \frac{\mathbb{V}(\theta|\mathbf{y})}{m}\right) \quad \text{for large } m$$

Monte Carlo sampling - convergence



Gibbs sampling

- **Sampling from multivariate distributions**, $p(X_1, \dots, X_p)$.
- Typically a posterior distribution: $p(\theta_1, \dots, \theta_p | \mathbf{y})$.
- Needed: easily sampled **full conditional posterior distributions**:
 - ▶ $p(\theta_1 | \theta_2, \theta_3, \dots, \theta_p, \mathbf{y})$
 - ▶ $p(\theta_2 | \theta_1, \theta_3, \dots, \theta_p, \mathbf{y})$
 - ▶ \vdots
 - ▶ $p(\theta_p | \theta_1, \theta_2, \dots, \theta_{p-1}, \mathbf{y})$

The Gibbs sampling algorithm

Gibbs sampling

Input: initial values $\theta_2^{(0)}, \dots, \theta_p^{(0)}$
number of posterior draws m .

for i in $1:m$ **do**

$$\theta_1 \sim p(\theta_1 \mid \theta_2^{(i-1)}, \theta_3^{(i-1)}, \dots, \theta_p^{(i-1)}, \mathbf{y})$$

$$\theta_2 \sim p(\theta_2 \mid \theta_1^{(i)}, \theta_3^{(i-1)}, \dots, \theta_p^{(i-1)}, \mathbf{y})$$

$$\vdots$$

$$\theta_p \sim p(\theta_p \mid \theta_1^{(i)}, \theta_2^{(i)}, \dots, \theta_{p-1}^{(i)}, \mathbf{y})$$

end

Output: m autocorrelated draws for $\boldsymbol{\theta} = (\theta_1, \dots, \theta_p)^\top$
that converge in distribution to the joint
posterior $p(\theta_1, \dots, \theta_p \mid \mathbf{y})$.

Dependent draws for Gibbs are less efficient

- $\theta^{(1)}, \dots, \theta^{(m)}$ **converges in distribution** to posterior $p(\theta|\mathbf{y})$.
- **Dependent draws** \rightarrow **less efficient** than iid sampling.
- **IID samples:**

$$\text{Var}(\bar{\theta}) = \frac{\sigma^2}{m}, \quad \text{where } \sigma^2 = \mathbb{V}(\theta|\mathbf{y})$$

- **Autocorrelated samples:**

$$\text{Var}(\bar{\theta}) = \frac{\sigma^2}{m} \left(1 + 2 \sum_{k=1}^{\infty} \rho_k \right)$$

where ρ_k is the autocorrelation at lag k .

- **Inefficiency factor:**

$$\text{IF} = 1 + 2 \sum_{k=1}^{\infty} \rho_k \approx 1 + 2 \sum_{k=1}^K \rho_k$$

- **Effective sample size (ESS):** $\frac{m}{\text{IF}}$.

Gibbs sampling bivariate normal

■ Joint distribution

$$\begin{pmatrix} \theta_1 \\ \theta_2 \end{pmatrix} \sim N_2 \left[\begin{pmatrix} \mu_1 \\ \mu_2 \end{pmatrix}, \begin{pmatrix} \sigma_1^2 & \rho\sigma_1\sigma_2 \\ \rho\sigma_1\sigma_2 & \sigma_2^2 \end{pmatrix} \right]$$

Gibbs sampling from a bivariate normal

Input: initial value $\theta_2^{(0)}$
number of posterior draws m .

for i in $1:m$ **do**

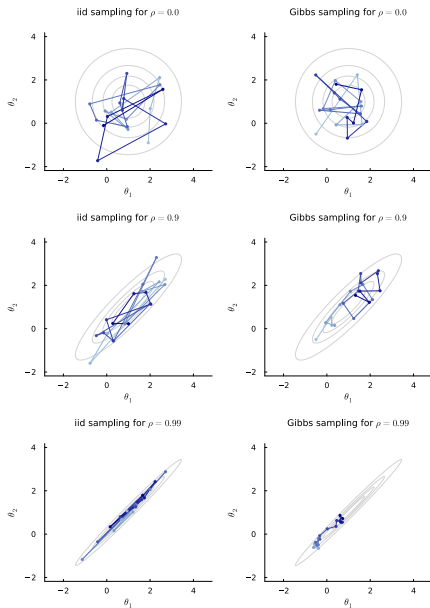
$$\begin{array}{|l} \theta_1^{(i)} \mid \theta_2 \sim N(\mu_1 + \rho \frac{\sigma_1}{\sigma_2} (\theta_2^{(i-1)} - \mu_2), \sigma_1^2(1-\rho)^2) \\ \theta_2^{(i)} \mid \theta_1 \sim N(\mu_2 + \rho \frac{\sigma_2}{\sigma_1} (\theta_1^{(i)} - \mu_1), \sigma_2^2(1-\rho)^2) \end{array}$$

end







Output: m autocorrelated draws for $\theta = (\theta_1, \theta_2)^\top$ that converge in distribution to the bivariate normal distribution $\theta \sim N(\mu, \Sigma)$, where $\mu = (\mu_1, \mu_2)^\top$ and

$$\Sigma = \begin{pmatrix} \sigma_1^2 & \rho\sigma_1\sigma_2 \\ \rho\sigma_1\sigma_2 & \sigma_2^2 \end{pmatrix}$$


Direct sampling vs Gibbs sampling



Direct vs Gibbs sampling, bivariate normal $\rho = 0.9$



Gibbs is inefficient when parameters are correlated



Bayesian learning of ridge regularization parameter

- Cross-validation is often used to determine λ .
- Bayesian: λ is **unknown** \Rightarrow **use a prior** for λ .
- $\lambda^{-1} \sim \text{Inv-}\chi^2(\omega_0, \psi_0^2)$. The user specifies ω_0 and ψ_0^2 .
- Joint posterior
$$p(\beta, \sigma^2, \lambda | \mathbf{y}, \mathbf{X})$$
- Marginal posterior λ .
- Gibbs sampling

Gibbs sampling for ridge regularization parameter

Gibbs sampling linear regression - L2 regularization prior

The posterior for the linear regression model

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \varepsilon, \varepsilon \sim N(\mathbf{0}, \sigma^2 I_n), \quad (11.16)$$

with hierarchical L2 regularization prior

$$\begin{aligned}\boldsymbol{\beta} | \sigma^2, \lambda &\sim N(\mathbf{0}, (\sigma^2 / \lambda) I_p) \\ \sigma^2 &\sim \text{Inv-}\chi^2(\tau_0^2, \nu_0) \\ \lambda^{-1} &\sim \text{Inv-}\chi^2(\omega_0, \psi_0^2).\end{aligned}$$

can be sampled by a two-block Gibbs sampler:

$$\begin{aligned}\text{Block1 : } \boldsymbol{\beta} | \sigma^2, \lambda, \mathbf{y} &\sim N(\hat{\boldsymbol{\beta}}_{L_2}, \sigma^2 (\mathbf{X}^\top \mathbf{X} + \lambda I_p)^{-1}) \\ \sigma^2 | \lambda, \mathbf{y} &\sim \text{Inv-}\chi^2(\tau_n^2, \nu_n)\end{aligned}$$

$$\text{Block2 : } \lambda^{-1} | \boldsymbol{\beta}, \sigma^2, \mathbf{y} \sim \text{Inv-}\chi^2(\omega_n, \psi_n^2),$$

Gibbs sampling

- Gibbs sampling can be the best approach:
 - ▶ when **correlated parameters can be sampled as a block**
 - ▶ when sampling of the large blocks is fast
- Other samplers (e.g. **Metropolis-Hastings**) **within Gibbs**.
- **Data augmentation** can make Gibbs more applicable:
 - ▶ Probit regression
 - ▶ Logistic regression
 - ▶ Horseshoe-regularized regression and classification
 - ▶ Mixture models

Data augmentation - Probit regression

■ Probit regression:

$$\Pr(y_i = 1 \mid \mathbf{x}_i) = \Phi(\mathbf{x}_i^\top \boldsymbol{\beta})$$

■ Random utility formulation:

$$\begin{aligned} u_i &\sim N(\mathbf{x}_i^\top \boldsymbol{\beta}, 1) \\ y_i &= \begin{cases} 1 & \text{if } u_i > 0 \\ 0 & \text{if } u_i \leq 0 \end{cases} \end{aligned}$$

■ Gibbs sampling samples from $p(\boldsymbol{\beta}, u_1, \dots, u_n \mid \mathbf{y}, \mathbf{X})$:

- ▶ Sample $\boldsymbol{\beta} \mid \mathbf{u}, \mathbf{y}, \mathbf{X}$ using linear regression update
- ▶ Sample each $u_i \mid \boldsymbol{\beta}, \mathbf{y}, \mathbf{X}$ using truncated normal

Random walk Metropolis algorithm

■ **Initialize** $\theta^{(0)}$ and iterate for $i = 1, 2, \dots$


1 **Sample proposal:** $\theta_p | \theta^{(i-1)} \sim N(\theta^{(i-1)}, c \cdot \Sigma)$

2 Compute the **acceptance probability**

$$\alpha = \min \left(1, \frac{p(\mathbf{y} | \theta_p) p(\theta_p)}{p(\theta^{(i-1)} | \mathbf{y})} \right)$$

3 With probability α set $\theta^{(i)} = \theta_p$ and $\theta^{(i)} = \theta^{(i-1)}$ otherwise.

00 Random walk Metropolis



Number of iterations



250

Standard deviation of proposals



1

This parameter controls how far the sampler can jump in each iteration.

Random walk Metropolis, cont.

- Common choices of Σ in proposal $N(\theta^{(i-1)}, c \cdot \Sigma)$:
 - ▶ $\Sigma = I$ (proposes 'off the cigar')
 - ▶ $\Sigma = J_y^{-1}(\hat{\theta})$ (propose 'along the cigar')
 - ▶ **Adaptive**. Start with $\Sigma = I$. Update Σ from initial run.
- Set c so average acceptance probability is 25-30%.
- **Good proposal**:
 - ▶ **Easy to sample**
 - ▶ **Easy to compute** α
 - ▶ Proposals should take reasonably **large steps** in θ -space
 - ▶ Proposals should **not be reject too often**.

The Metropolis-Hastings algorithm

- Generalization when the proposal density is not symmetric.

- Initialize $\theta^{(0)}$ and iterate for $i = 1, 2, \dots$

- 1 **Sample proposal:** $\theta_p \sim q(\cdot | \theta^{(i-1)})$




- 2 Compute the **acceptance probability**

$$\alpha = \min \left(1, \frac{p(\mathbf{y} | \theta_p) p(\theta_p)}{p(\mathbf{y} | \theta^{(i-1)}) p(\theta^{(i-1)})} \frac{q(\theta^{(i-1)} | \theta_p)}{q(\theta_p | \theta^{(i-1)})} \right)$$

- 3 With probability α set $\theta^{(i)} = \theta_p$ and otherwise $\theta^{(i)} = \theta^{(i-1)}$


Burn-in and convergence

- How long **burn-in**?
- **How long to sample** after burn-in?
- **Thinning**? Keeping every h draw reduces autocorrelation.
- **Convergence diagnostics**
 - ▶ Raw plots of simulated sequences (trajectories)
 - ▶ CUSUM plots
 - ▶ Check multiple runs with different initial values.
 - ▶ Potential scale reduction factor, \hat{R} .




Burn-in and convergence


nSim = 500




nSim = 1000



nSim = 1500



nSim = 2000



Hamiltonian Monte Carlo

- When $\theta = (\theta_1, \dots, \theta_p)^\top$ is **high-dimensional**, $p(\theta|\mathbf{y})$ usually located in some subregion of \mathbb{R}^p with complicated geometry.
- MH: hard to find good proposal distribution $q(\cdot|\theta^{(i-1)})$.
- MH: use very small step sizes otherwise too many rejections.
- **Hamiltonian Monte Carlo (HMC):**
 - ▶ distant proposals **and**
 - ▶ high acceptance probabilities.
- **HMC** is a **Metropolis-Hastings** + **proposal tailored to the posterior distribution** using ideas from Physics.

Hamiltonian Monte Carlo

- Physics: **Hamiltonian** system $H(\boldsymbol{\theta}, \boldsymbol{\phi}) = U(\boldsymbol{\theta}) + K(\boldsymbol{\phi})$, where U is the **potential energy** and K is the **kinetic energy**.
- **Hamiltonian Dynamics**

$$\begin{aligned}\frac{d\theta_i}{dt} &= \frac{\partial H}{\partial \phi_i} = \frac{\partial K}{\partial \phi_i}, \\ \frac{d\phi_i}{dt} &= -\frac{\partial H}{\partial \theta_i} = -\frac{\partial U}{\partial \theta_i}\end{aligned}$$

- Hockey puck sliding over a friction-less surface: [illustration](#).
- **Posterior sampling**: $U(\boldsymbol{\theta}) = -\log[p(\boldsymbol{\theta})p(\mathbf{y}|\boldsymbol{\theta})]$.
- Momentum: $\boldsymbol{\phi} \sim N(\mathbf{0}, \mathbf{M})$ where \mathbf{M} is the mass matrix and

$$K(\boldsymbol{\phi}) = -\log[p(\boldsymbol{\phi})] = \frac{1}{2}\boldsymbol{\phi}^\top \mathbf{M}^{-1}\boldsymbol{\phi} + \text{const}$$

- If we could propose $\boldsymbol{\theta}$ in continuous time (spoiler: we can't), the acceptance probability would be one.

Hamiltonian Monte Carlo

■ Hamiltonian Dynamics

$$\begin{aligned}\frac{d\theta_i}{dt} &= [\mathbf{M}^{-1}\boldsymbol{\phi}]_i, \\ \frac{d\phi_i}{dt} &= \frac{\partial \log p(\boldsymbol{\theta}|\mathbf{y})}{\partial \theta_i}\end{aligned}$$


approximated using L steps with the **leapfrog algorithm**

$$\begin{aligned}\phi_i\left(t + \frac{\varepsilon}{2}\right) &= \phi_i(t) + \frac{\varepsilon}{2} \frac{\partial \log p(\boldsymbol{\theta}|\mathbf{y})}{\partial \theta_i} \Big|_{\theta(t)} \\ \theta_i(t + \varepsilon) &= \theta_i(t) + \varepsilon \mathbf{M}^{-1} \phi_i\left(t + \frac{\varepsilon}{2}\right), \\ \phi_i(t + \varepsilon) &= \phi_i\left(t + \frac{\varepsilon}{2}\right) + \frac{\varepsilon}{2} \frac{\partial \log p(\boldsymbol{\theta}|\mathbf{y})}{\partial \theta_i} \Big|_{\theta(t+\varepsilon)},\end{aligned}$$

where ε is the **step size**.

■ **Discretization** \Rightarrow acceptance probability drops with ε .

Leapfrog algorithm



Step size

0.2

How far the integrator moves at each step

Trajectory length

33

Number of steps to take in the trajectory

The Hamiltonian Monte Carlo algorithm


■ Initialize $\theta^{(0)}$ and iterate for $i = 1, 2, \dots$

- 1 Sample the starting **momentum** $\phi_s \sim N(0, \mathbf{M})$
- 2 Simulate new values for (θ_p, ϕ_p) by iterating the **leapfrog algorithm** L times with step size ε , starting in $(\theta^{(i-1)}, \phi_s)$.
- 3 Compute the **acceptance probability**

$$\alpha = \min \left(1, \frac{p(\mathbf{y}|\theta_p)p(\theta_p)}{p(\mathbf{y}|\theta^{(i-1)})p(\theta^{(i-1)})} \frac{p(\phi_p)}{p(\phi_s)} \right)$$

- 4 With probability α set $\theta^{(i)} = \theta_p$ and $\theta^{(i)} = \theta^{(i-1)}$ otherwise.

Hamiltonian Monte Carlo



The Hamiltonian Monte Carlo algorithm

■ Initialize $\theta^{(0)}$ and iterate for $i = 1, 2, \dots$

- 1 Sample the starting **momentum** $\phi_s \sim N(0, \mathbf{M})$
- 2 Simulate new values for (θ_p, ϕ_p) by iterating the **leapfrog algorithm** L times with step size ε , starting in $(\theta^{(i-1)}, \phi_s)$.
- 3 Compute the **acceptance probability**


$$\alpha = \min \left(1, \frac{p(\mathbf{y}|\theta_p)p(\theta_p)}{p(\mathbf{y}|\theta^{(i-1)})p(\theta^{(i-1)})} \frac{p(\phi_p)}{p(\phi_s)} \right)$$

- 4 With probability α set $\theta^{(i)} = \theta_p$ and $\theta^{(i)} = \theta^{(i-1)}$ otherwise.

Tuning Hamiltonian Monte Carlo


- HMC is very efficient, but **needs careful tuning** to work.
- **Tuning parameters:**
 - ▶ **stepsize** ε ,
 - ▶ **number of leapfrog** iterations L and
 - ▶ **mass matrix** M . (hello $J_x^{-1}(\hat{\theta})$, my old friend)
- **No U-turn** sampler:
 - ▶ **Warm-up** to determine ε and L to get good acceptance rate.
 - ▶ Avoids U-turns in the Hamiltonian proposals.
- Drawbacks of HMC:
 - ▶ Need to **evaluate gradient of log posterior** many times during Hamiltonian iterations. Costly! (Subsampling HMC).
 - ▶ Difficulty with **multimodality** (true for most algorithms).
 - ▶ Standard HMC cannot handle **discrete parameters**. Mixture example. Some recent progress.

Hamiltonian Monte Carlo on multimodal posterior




Comparing algorithms for bivariate normal


Gibbs




RWM - Identity M



HMC - Diagonal M



HMC - Full M



Comparing algorithms for bivariate normal

