

Workshop: Intro to Bayesian Learning

Lecture 6 - Implementing Bayesian Learning with Probabilistic Programming

Mattias Villani

**Department of Statistics
Stockholm University**



Overview


- Probabilistic programming
- Turing
- Stan

Probabilistic programming languages for Bayes

- **Stan** is a probabilistic programming language for Bayes based on HMC.
- C++ using the R package `rstan`. Bindings from Python.



- **Turing.jl** is a probabilistic programming language in Julia.
- Written in Julia, which is fast natively.



HMC sampling for iid normal model in rstan

```
library(rstan)

# Define the Stan model
stanModelNormal = '
// The input data is a vector y of length N.
data {
  // data
  int<lower=0> N;
  vector[N] y;
  // prior
  real mu0;
  real<lower=0> kappa0;
  real<lower=0> nu0;
  real<lower=0> sigma20;
}

// The parameters in the model
parameters {
  real theta;
  real<lower=0> sigma2;
}

model {
  sigma2 ~ scaled_inv_chi_square(nu0, sqrt(sigma20));
  theta ~ normal(mu0, sqrt(sigma2/kappa0));
  y ~ normal(theta, sqrt(sigma2));
}
'

# Set up the observed data
data <- list(N = 5, y = c(15.77, 20.5, 8.26, 14.37, 21.09))

# Set up the prior
prior <- list(mu0 = 20, kappa0 = 1, nu0 = 5, sigma20 = 5^2)

# Sample from posterior using HMC
fit <- stan(model_code = stanModelNormal, data = c(data,prior), iter = 10000 )
```

HMC sampling for iid normal model in Turing.jl

```
using Turing

ScaledInverseChiSq(v,τ²) = InverseGamma(v/2,v*τ²/2) # Scaled Inv- $\chi^2$  distribution

# Setting up the Turing model:
@model function iidnormal(x, μ₀, κ₀, v₀, σ²₀)
    σ² ~ ScaledInverseChiSq(v₀, σ²₀)
    θ ~ Normal(μ₀, √(σ²/κ₀)) # prior
    n = length(x) # number of observations
    for i in 1:n
        x[i] ~ Normal(θ, √σ²) # model
    end
end

# Set up the observed data
x = [15.77,20.5,8.26,14.37,21.09]

# Set up the prior
μ₀ = 20; κ₀ = 1; v₀ = 5; σ²₀ = 5^2

# Settings of the Hamiltonian Monte Carlo (HMC) sampler.
α = 0.8
postdraws = sample(iidnormal(x, μ₀, κ₀, v₀, σ²₀), NUTS(α), 10000, discard_initial = 1000)
```

Modeling the number of bidders in eBay auctions

variable	description	data type	original range
nbids	number of bids	counts	[0, 12]
bookvalue	coin's book value	continuous	[7.5, 399.5]
startprice	seller's reservation price / book value	continuous	[0, 1.702]
minblemish	minor blemish	binary	[0, 1]
majblemish	major blemish	binary	[0, 1]
negfeedback	large negative feedback score	binary	[0, 1]
powerseller	large quantity seller	binary	[0, 1]
verified	verified seller on ebay	binary	[0, 1]
sealed	unopened package	binary	[0, 1]

■ Poisson regression

$$y_i | \mathbf{x}_i \sim \text{Poisson}(\lambda_i)$$

$$\lambda_i = \exp(\mathbf{x}_i^\top \boldsymbol{\beta})$$

HMC sampling for Poisson regression in Turing.jl

```
using Turing

# Setting up the poisson regression model
@model function poissonReg(y, X, τ)
    p = size(X,2)
    β ~ filldist(Normal(0, τ), p) # all  $\beta_j$  are iid Normal(0, τ)
    λ = exp.(X*β)
    n = length(y)
    for i in 1:n
        y[i] ~ Poisson(λ[i])
    end
end

# HMC sampling from posterior of β
τ = 10 # Prior standard deviation
α = 0.70 # target acceptance probability in NUTS sampler
model = poissonReg(y, X, τ)
chain = sample(model, Turing.NUTS(α), 10000, discard_initial = 1000)
```

■ [Poisson regression](#) in rstan.

... or TuringGLM.jl with R's formula syntax

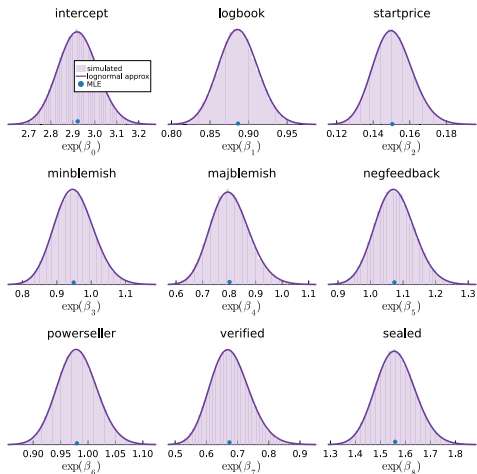
```
# Using TuringGLM.jl
using TuringGLM
fm = @formula(nbids ~ logbook + startprice + minblemish +
|      majblemish + negfeedback + powerseller + verified + sealed)
model = turing_model(fm, ebay_df; model = Poisson)
chain = sample(model, NUTS(), 10000)
```

- Inspired by the [brms](#) package in R.

Marginal posteriors

■ Multiplicative model

$$E(y|\mathbf{x}) = \exp(\beta_0 + \beta_1 x_1 + \beta_2 x_2) = \exp(\beta_0) \exp(\beta_1)^{x_1} \exp(\beta_2)^{x_2}$$



Negative binomial regression in Turing.jl


■ Negative binomial regression

$$y_i | \mathbf{x}_i \sim \text{NegBinomial} \left(\psi, p = \frac{\psi}{\psi + \lambda_i} \right), \quad \lambda_i = \exp(\mathbf{x}_i^\top \boldsymbol{\beta})$$

■ Mean is still λ_i , but variance is larger: $\text{Var}(y_i) = \lambda_i(1 + \lambda_i/\psi)$.

■ As $\psi \rightarrow \infty$ we get Poisson again.

```
# Negative binomial regression
@model function negbinomialReg(y, X, τ, μ₀, σ₀)
    p = size(X,2)
    β ~ filldist(Normal(0, τ), p)
    λ = exp.(X*β)
    ψ ~ LogNormal(μ₀, σ₀)
    n = length(y)
    for i in 1:n
        y[i] ~ NegativeBinomial(ψ, ψ/(ψ + λ[i]))
    end
end
```



Regression with horseshoe in Turing.jl

■ Horseshoe prior

$$\beta_j | \lambda_j^2, \tau^2 \stackrel{\text{ind}}{\sim} N(0, \sigma^2 \tau^2 \lambda_j^2) \quad \lambda_j \sim C^+(0, 1) \quad \tau \sim C^+(0, 1)$$

```
# Define the half-Cauchy distribution
```

```
halfCauchy = truncated(Cauchy(0, 1); lower=0)
```

```
@model function BayesLinRegHS_turing(y, X, v_o, σ²_o, stdβ_o = 100)
```

```
    p = size(X, 2) # X should not include intercept
```

```
    τ ~ halfCauchy
```

```
    λ ~ filldist(halfCauchy, p)
```

```
    σ² ~ ScaledInverseChiSq(v_o, σ²_o)
```

```
    β_o ~ Normal(0, stdβ_o)
```

```
    β ~ MvNormal(zeros(p), Diagonal((λ .* τ).^2 .* σ²))
```

```
    y ~ MvNormal(β_o .+ X*β, σ²*I)
```

```
end
```

```
# Prior hyperparameters
```

```
v_o = 0.01; σ²_o = 1.0
```

```
# Run HMC to sample from the posterior
```

```
@time hmc_draws = sample(BayesLinRegHS_turing(y, X[:,2:end], v_o, σ²_o),
```

```
    NUTS(), nDraws, n_adapts = 1000, n_chains = 1);
```

```
# Plot results
```

```
histogram(hmc_draws[:,τ], bins = 100)
```

Regression with horseshoe using Gibbs sampling

```
for i ∈ 1:nSim

    # Compute things needed for sampling  $\beta, \sigma \mid \lambda, \tau, y, X$ 
    invLambda = diagm(1 ./ (lambda.^2))
    Omega_o = [1/(stdbeta_o.^2) zeros(p,1) (1/tau.^2)*invLambda]
    Omega_n = Symmetric(XX + Omega_o)
    invOmega_n = inv(Omega_n)
    mu_n = Omega_n \ (XX * beta_hat)
    sigma2_n = (v_o * sigma2_o + (y - X * beta_hat)' * (y - X * beta_hat) + (mu_n - beta_hat)' * XX * (mu_n - beta_hat) + mu_n' * Omega_o * mu_n) / v_n

    # Simulate from  $p(\sigma^2 \mid \psi^2, y, X)$ 
    sigma2 = rand(ScaledInverseChiSq(v_n, sigma2_n))
    sigma2sim[i] = sigma2
    sigma = sqrt(sigma2)

    # Simulate from  $p(\beta \mid \psi^2, \sigma^2, y, X)$ 
    beta = rand(MvNormal(mu_n, sigma2 * invOmega_n))
    beta_sim[i,:] = beta'

    # Simulate from  $p(\lambda \mid \tau, \beta, \sigma^2, y, X)$ 
    v = rand.( ScaledInverseChiSq.(2, 1 + 1 ./ (lambda.^2)) )
    lambda = sqrt.(rand.(
        ScaledInverseChiSq.(2, 1 ./ v .+ 0.5*(beta[intercept+1:end]/(sigma*tau)).^2 )
    ))

    # Simulate from  $p(\tau \mid \lambda, \beta, \sigma^2, y, X)$ 
    if estimate_tau
        xi = rand( ScaledInverseChiSq.(2, 1 + 1/(tau.^2)) )
        tau = sqrt(rand(
            ScaledInverseChiSq(p+1, (2/xi + sum((beta[intercept+1:end] ./ (sigma*lambda)).^2 ))/(p+1) )
        ))
    end
    tau_sim[i] = tau

end
```

end